

Data 621 - Homework 1

Group 4 Layla Quinones, Ian Castello, Dmitriy Burtsev & Esteban Aramayo

September 26, 2021

About the Data

The data set consists of 2,276 records and 17 different variables, with each observation corresponding to a baseball teams performance in a single year. The time horizon of these data are from 1871, the same year as the first recorded professional baseball game through 2006.

General objective

Through linear regression, train the data to predict the number of wins.

Challenges Right Off the Bat (so to speak)

The data set covers a very large time period. The rules and play style of baseball have changed a great deal from the late 19th-century. The season year of the team would be an important factor in improving these models. Additionally, certain clubs have bucked trends in winning or losing despite these metrics. The Boston Red Sox and Chicago Cubs had very long dry spells, even with good team statistics.

Data Exploration

Our data is stored for easy reference among the team on GitHub. With 2,276 team observations and 17 variables. Of those, 15 are features, 1 is an index, and the remaining is our target variable for number of wins. Right away, we already know that missing values will need to be accounted for in all of this features.

```
## 'data.frame': 2276 obs. of 17 variables:
## $ INDEX : int 1 2 3 4 5 6 7 8 11 12 ...
## $ TARGET_WINS : int 39 70 86 70 82 75 80 85 86 76 ...
## $ TEAM_BATTING_H : int 1445 1339 1377 1387 1297 1279 1244 1273 1391 1271 ...
## $ TEAM_BATTING_2B : int 194 219 232 209 186 200 179 171 197 213 ...
## $ TEAM_BATTING_3B : int 39 22 35 38 27 36 54 37 40 18 ...
## $ TEAM_BATTING_HR : int 13 190 137 96 102 92 122 115 114 96 ...
## $ TEAM_BATTING_BB : int 143 685 602 451 472 443 525 456 447 441 ...
## $ TEAM_BATTING_SO : int 842 1075 917 922 920 973 1062 1027 922 827 ...
## $ TEAM_BASERUN_SB : int NA 37 46 43 49 107 80 40 69 72 ...
## $ TEAM_BASERUN_CS : int NA 28 27 30 39 59 54 36 27 34 ...
## $ TEAM_BATTING_HBP: int NA NA NA NA NA NA NA NA NA NA ...
## $ TEAM_PITCHING_H : int 9364 1347 1377 1396 1297 1279 1244 1281 1391 1271 ...
## $ TEAM_PITCHING_HR: int 84 191 137 97 102 92 122 116 114 96 ...
```

```
## $ TEAM_PITCHING_BB: int 927 689 602 454 472 443 525 459 447 441 ...
## $ TEAM_PITCHING_SO: int 5456 1082 917 928 920 973 1062 1033 922 827 ...
## $ TEAM_FIELDING_E : int 1011 193 175 164 138 123 136 112 127 131 ...
## $ TEAM_FIELDING_DP: int NA 155 153 156 168 149 186 136 169 159 ...
```










Summary Statistics

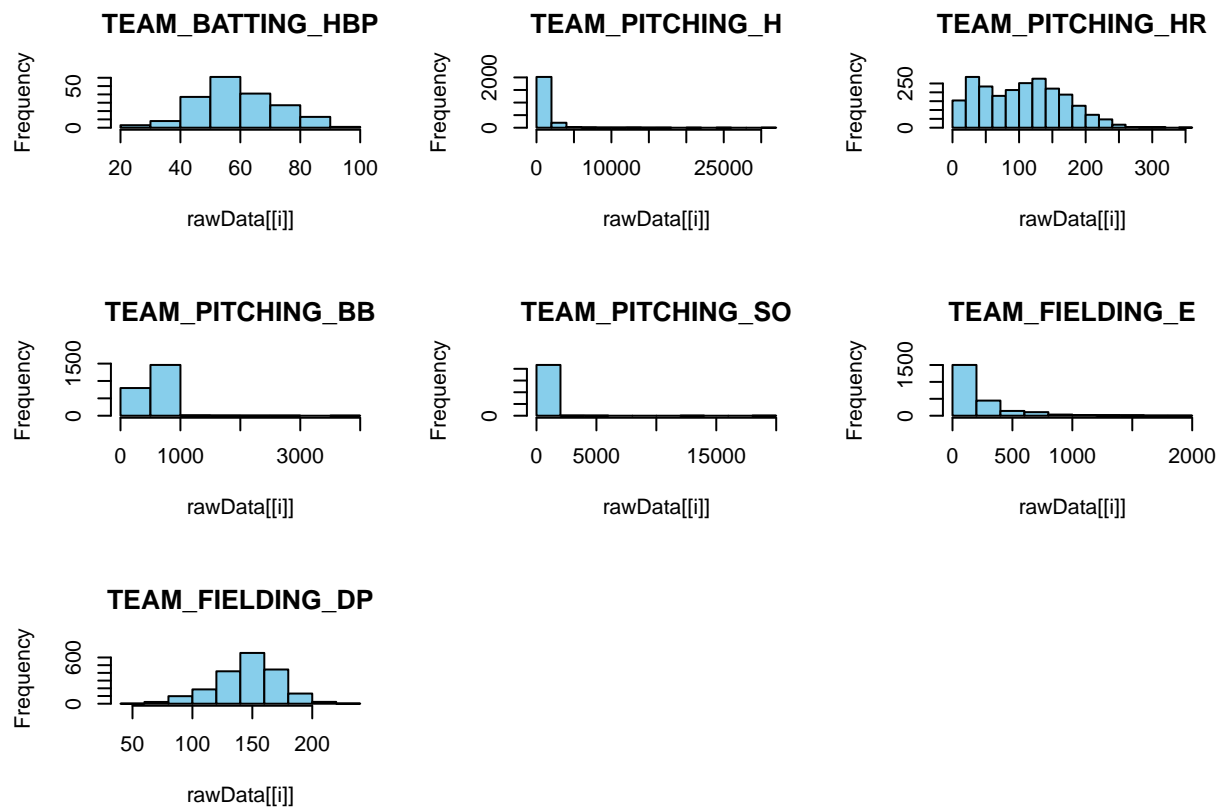
For each of the variables, these summary statistics provide a nice overview of each feature, its variation, and paths for potential transformations later on for model construction. Of note are the normally distributed variables, like our target variable for wins, base hits by batters, doubles by batters, walks by batters, and batters hit by pitches. The more skewed features include hits allowed, strike outs by pitchers (a very difficult thing to do consistently), and team fielding errors. Once again, we can observe the extent of the N/As and outliers that we'll have to account for.

```
##      INDEX      TARGET_WINS      TEAM_BATTING_H TEAM_BATTING_2B
## Min.   : 1.0    Min.   : 0.00    Min.   : 891    Min.   : 69.0
## 1st Qu.: 630.8  1st Qu.: 71.00    1st Qu.:1383    1st Qu.:208.0
## Median :1270.5  Median : 82.00    Median :1454    Median :238.0
## Mean   :1268.5  Mean   : 80.79    Mean   :1469    Mean   :241.2
## 3rd Qu.:1915.5  3rd Qu.: 92.00    3rd Qu.:1537    3rd Qu.:273.0
## Max.   :2535.0  Max.   :146.00    Max.   :2554    Max.   :458.0
##
## TEAM_BATTING_3B TEAM_BATTING_HR TEAM_BATTING_BB TEAM_BATTING_SO
## Min.   : 0.00    Min.   : 0.00    Min.   : 0.0    Min.   : 0.0
## 1st Qu.: 34.00    1st Qu.: 42.00    1st Qu.:451.0    1st Qu.: 548.0
## Median : 47.00    Median :102.00    Median :512.0    Median : 750.0
## Mean   : 55.25    Mean   : 99.61    Mean   :501.6    Mean   : 735.6
## 3rd Qu.: 72.00    3rd Qu.:147.00    3rd Qu.:580.0    3rd Qu.: 930.0
## Max.   :223.00    Max.   :264.00    Max.   :878.0    Max.   :1399.0
##                                     NA's   :102
## TEAM_BASERUN_SB TEAM_BASERUN_CS TEAM_BATTING_HBP TEAM_PITCHING_H
## Min.   : 0.0    Min.   : 0.0    Min.   :29.00    Min.   : 1137
## 1st Qu.: 66.0    1st Qu.: 38.0    1st Qu.:50.50    1st Qu.: 1419
## Median :101.0    Median : 49.0    Median :58.00    Median : 1518
## Mean   :124.8    Mean   : 52.8    Mean   :59.36    Mean   : 1779
## 3rd Qu.:156.0    3rd Qu.: 62.0    3rd Qu.:67.00    3rd Qu.: 1682
## Max.   :697.0    Max.   :201.0    Max.   :95.00    Max.   :30132
## NA's   :131     NA's   :772     NA's   :2085
## TEAM_PITCHING_HR TEAM_PITCHING_BB TEAM_PITCHING_SO TEAM_FIELDING_E
## Min.   : 0.0    Min.   : 0.0    Min.   : 0.0    Min.   : 65.0
## 1st Qu.: 50.0    1st Qu.: 476.0    1st Qu.: 615.0    1st Qu.: 127.0
## Median :107.0    Median : 536.5    Median : 813.5    Median : 159.0
## Mean   :105.7    Mean   : 553.0    Mean   : 817.7    Mean   : 246.5
## 3rd Qu.:150.0    3rd Qu.: 611.0    3rd Qu.: 968.0    3rd Qu.: 249.2
## Max.   :343.0    Max.   :3645.0    Max.   :19278.0    Max.   :1898.0
##                                     NA's   :102
## TEAM_FIELDING_DP
## Min.   : 52.0
## 1st Qu.:131.0
## Median :149.0
## Mean   :146.4
## 3rd Qu.:164.0
## Max.   :228.0
```

NA's :286

##		vars	n	mean	sd	median	trimmed	mad	min	max
##	INDEX	1	2276	1268.46	736.35	1270.5	1268.57	952.57	1	2535
##	TARGET_WINS	2	2276	80.79	15.75	82.0	81.31	14.83	0	146
##	TEAM_BATTING_H	3	2276	1469.27	144.59	1454.0	1459.04	114.16	891	2554
##	TEAM_BATTING_2B	4	2276	241.25	46.80	238.0	240.40	47.44	69	458
##	TEAM_BATTING_3B	5	2276	55.25	27.94	47.0	52.18	23.72	0	223
##	TEAM_BATTING_HR	6	2276	99.61	60.55	102.0	97.39	78.58	0	264
##	TEAM_BATTING_BB	7	2276	501.56	122.67	512.0	512.18	94.89	0	878
##	TEAM_BATTING_SO	8	2174	735.61	248.53	750.0	742.31	284.66	0	1399
##	TEAM_BASERUN_SB	9	2145	124.76	87.79	101.0	110.81	60.79	0	697
##	TEAM_BASERUN_CS	10	1504	52.80	22.96	49.0	50.36	17.79	0	201
##	TEAM_BATTING_HBP	11	191	59.36	12.97	58.0	58.86	11.86	29	95
##	TEAM_PITCHING_H	12	2276	1779.21	1406.84	1518.0	1555.90	174.95	1137	30132
##	TEAM_PITCHING_HR	13	2276	105.70	61.30	107.0	103.16	74.13	0	343
##	TEAM_PITCHING_BB	14	2276	553.01	166.36	536.5	542.62	98.59	0	3645
##	TEAM_PITCHING_SO	15	2174	817.73	553.09	813.5	796.93	257.23	0	19278
##	TEAM_FIELDING_E	16	2276	246.48	227.77	159.0	193.44	62.27	65	1898
##	TEAM_FIELDING_DP	17	1990	146.39	26.23	149.0	147.58	23.72	52	228
##		range	skew	kurtosis	se					
##	INDEX	2534	0.00	-1.22	15.43					
##	TARGET_WINS	146	-0.40	1.03	0.33					
##	TEAM_BATTING_H	1663	1.57	7.28	3.03					
##	TEAM_BATTING_2B	389	0.22	0.01	0.98					
##	TEAM_BATTING_3B	223	1.11	1.50	0.59					
##	TEAM_BATTING_HR	264	0.19	-0.96	1.27					
##	TEAM_BATTING_BB	878	-1.03	2.18	2.57					
##	TEAM_BATTING_SO	1399	-0.30	-0.32	5.33					
##	TEAM_BASERUN_SB	697	1.97	5.49	1.90					
##	TEAM_BASERUN_CS	201	1.98	7.62	0.59					
##	TEAM_BATTING_HBP	66	0.32	-0.11	0.94					
##	TEAM_PITCHING_H	28995	10.33	141.84	29.49					
##	TEAM_PITCHING_HR	343	0.29	-0.60	1.28					
##	TEAM_PITCHING_BB	3645	6.74	96.97	3.49					
##	TEAM_PITCHING_SO	19278	22.17	671.19	11.86					
##	TEAM_FIELDING_E	1833	2.99	10.97	4.77					
##	TEAM_FIELDING_DP	176	-0.39	0.18	0.59					

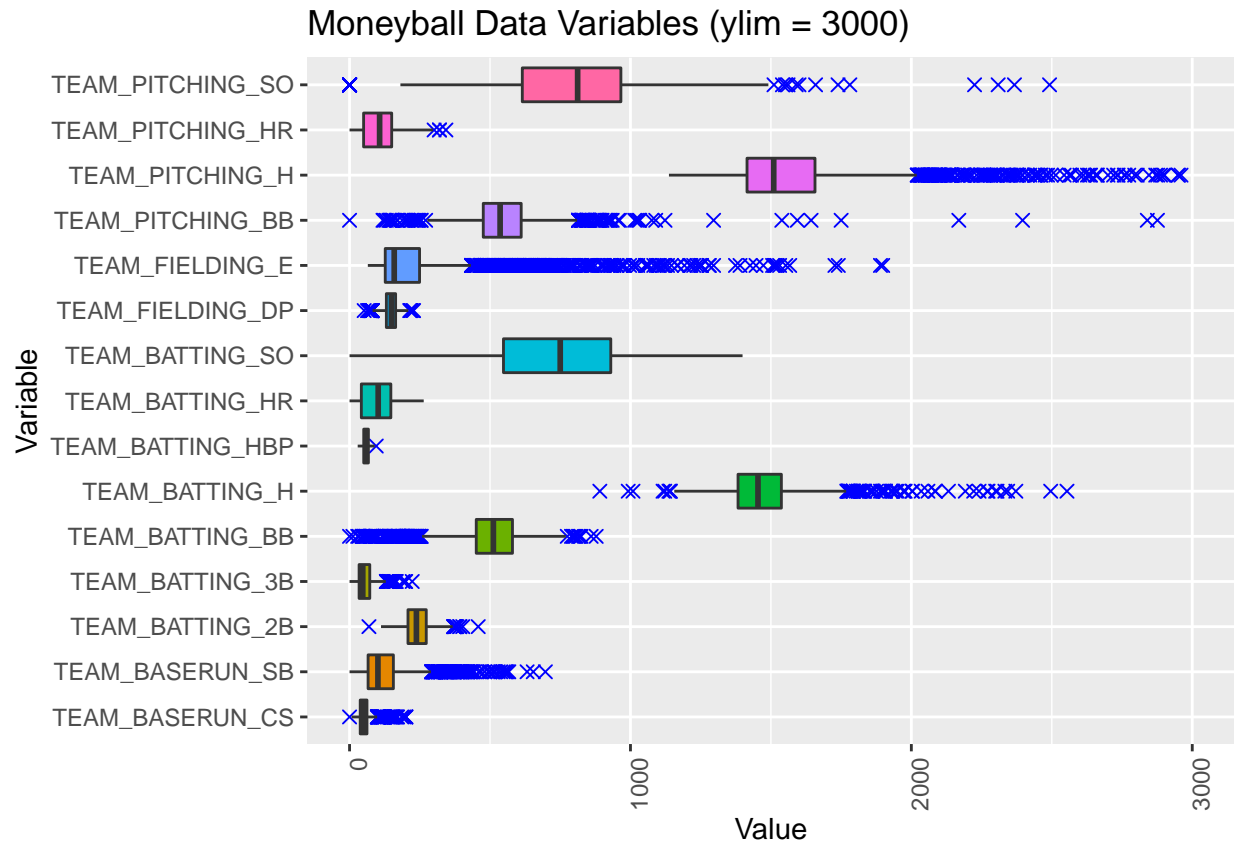




Box Plots

This box plot visualization gives us an idea of the outliers we have in each variable, but does not give us a good sense of the distribution. We can use the histograms above to interpret shape. From the box plots, we see that the variable `TEAM_PITCHING_H` has the greatest number of outliers. This may mean we throw that variable out altogether and not consider it in our models.

```
## Warning: Removed 3570 rows containing non-finite values (stat_boxplot).
```




Data Pre-Processing

Missing Data

Right away, there are a few variables with lots of missing data. Even with imputation, without at least 40% to 50% actual data, it would not be that informative to use in any model. Indeed, `TEAM_BATTING_HBP` has about 92% missing data and so is removed for the data set and will not be considered for the models. We will be using the histograms for each other variable to decide how and whether to impute during the pre-processing stage.


```
## Warning: It is deprecated to specify 'guide = FALSE' to remove a guide. Please
## use 'guide = "none"' instead.
```



```
## TEAM_BATTING_HBP TEAM_BASERUN_CS TEAM_FIELDING_DP TEAM_BASERUN_SB
## 1 0.9133772 0.339364 0.122807 0.05482456
## TEAM_BATTING_SO TEAM_PITCHING_SO
## 1 0.0433114 0.0433114
```

Correlation Matrix and Multicollinearity

It is important to check for features which may also be correlated. Simply, having multiple features relate to themselves can cause overfitting, reduced p values, and strange variances in the data. To avoid this, we exclude one or more of the variables. In this case, we see that `TEAM_BATTING_HR` and `TEAM_PITCHING_SO` are very intertwined, and we'll take care when constructing our models not to use both. We determined later that `TEAM_BATTING_HR` had very weak effects on the model and so selected that feature to be removed entirely.



```
## Compare row 6 and column 13 with corr 0.968
## Means: 0.374 vs 0.244 so flagging column 6
## All correlations <= 0.75
```

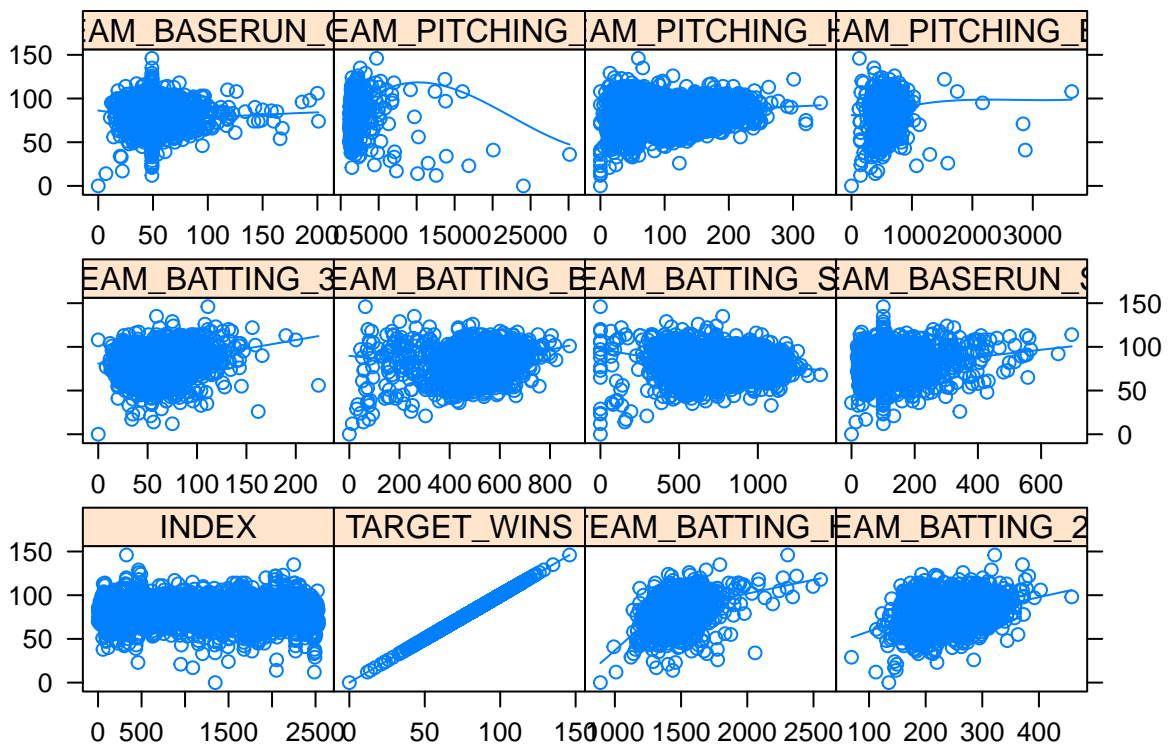
```
## [1] "TEAM_BATTING_HR"
```


Imputing Missing Data

For the other features with significant amounts missing data, we can impute using either the mean or median of the feature. In observing the histograms, if the shape is more skewed, we would seek to use the median. If the distribution appears more normal, we can use the mean (average).

Feature Plots

With the data cleaned and imputed, we can again review the features and begin selecting them for our models. The feature plots below summarize the potential effect of each feature on the target variable. Obviously, Target wins is basically a straight line since it is itself the target variable and a perfect line.



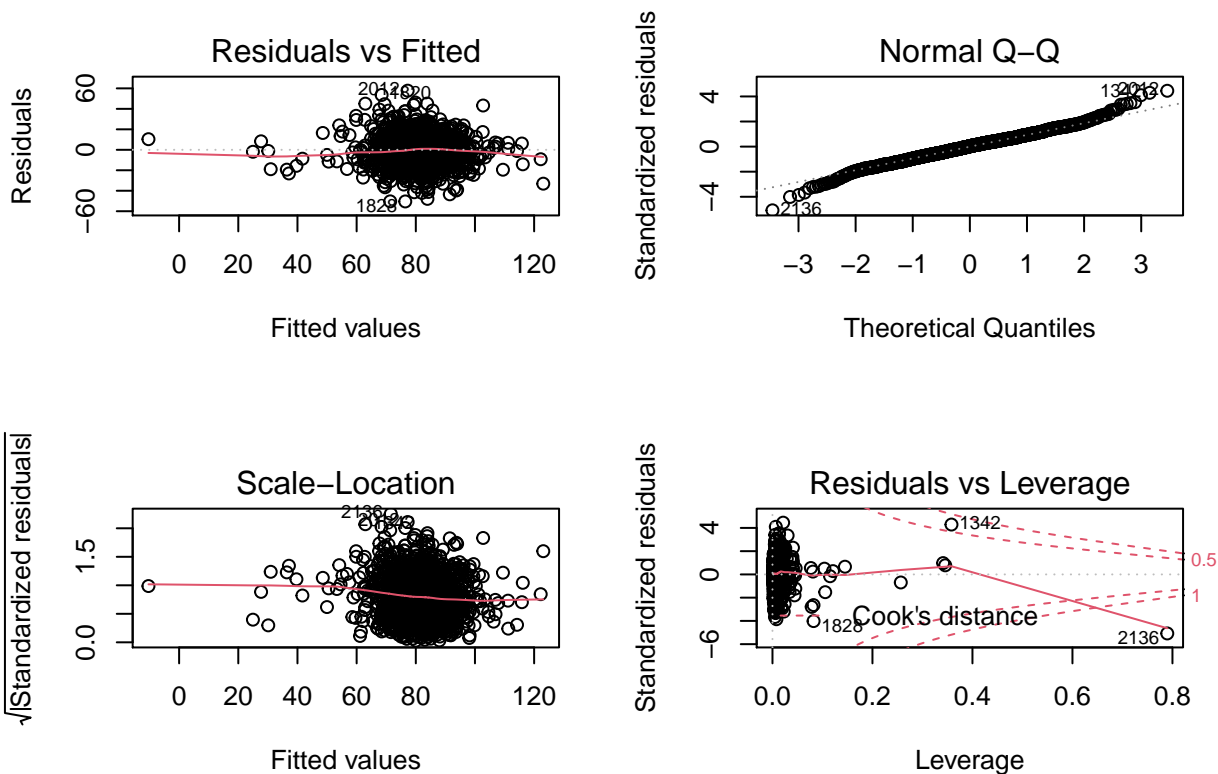


Build Models

Model 1: “The Kitchen Sink”

$$R^2 = 0.303$$

Our first model is the so-called kitchen sink approach, where all features are included. We get a pretty lousy R^2 value, even though the residuals appears to be quite random, if a bit clumped and the F-statistic indicated our model does say **something**.




```
##
## Call:
## lm(formula = TARGET_WINS ~ ., data = trainData)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -50.444  -8.364   0.404   8.140  57.744
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  23.4416666   6.0755834   3.858 0.000118 ***
## INDEX        -0.0006326   0.0004206  -1.504 0.132699
## TEAM_BATTING_H  0.0490246   0.0041222  11.893 < 2e-16 ***
## TEAM_BATTING_2B -0.0227145   0.0102269  -2.221 0.026471 *
## TEAM_BATTING_3B  0.0605813   0.0187539   3.230 0.001259 **
## TEAM_BATTING_BB  0.0150025   0.0059409   2.525 0.011645 *
## TEAM_BATTING_SO -0.0083507   0.0028737  -2.906 0.003706 **
## TEAM_BASERUN_SB  0.0252307   0.0048857   5.164 2.68e-07 ***
## TEAM_BASERUN_CS -0.0122584   0.0176581  -0.694 0.487640
## TEAM_PITCHING_H -0.0009030   0.0003871  -2.333 0.019775 *
## TEAM_PITCHING_HR  0.0588300   0.0097636   6.025 2.04e-09 ***
## TEAM_PITCHING_BB -0.0027049   0.0042154  -0.642 0.521174
## TEAM_PITCHING_SO  0.0036413   0.0010092   3.608 0.000317 ***
## TEAM_FIELDING_E -0.0183293   0.0027725  -6.611 5.00e-11 ***
## TEAM_FIELDING_DP -0.1154644   0.0144517  -7.990 2.39e-15 ***
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.11 on 1809 degrees of freedom
## Multiple R-squared:  0.3085, Adjusted R-squared:  0.3032
## F-statistic: 57.66 on 14 and 1809 DF,  p-value: < 2.2e-16
```

Model 2: Targeting Most Impactful Features

Reviewing the standard errors, and the correlation matrix, we what we think may be the most impactful metrics. Unfortunately, is has the opposite effect on our R^2 , which is reduced roughly 8%.

$$R^2 = 0.230$$



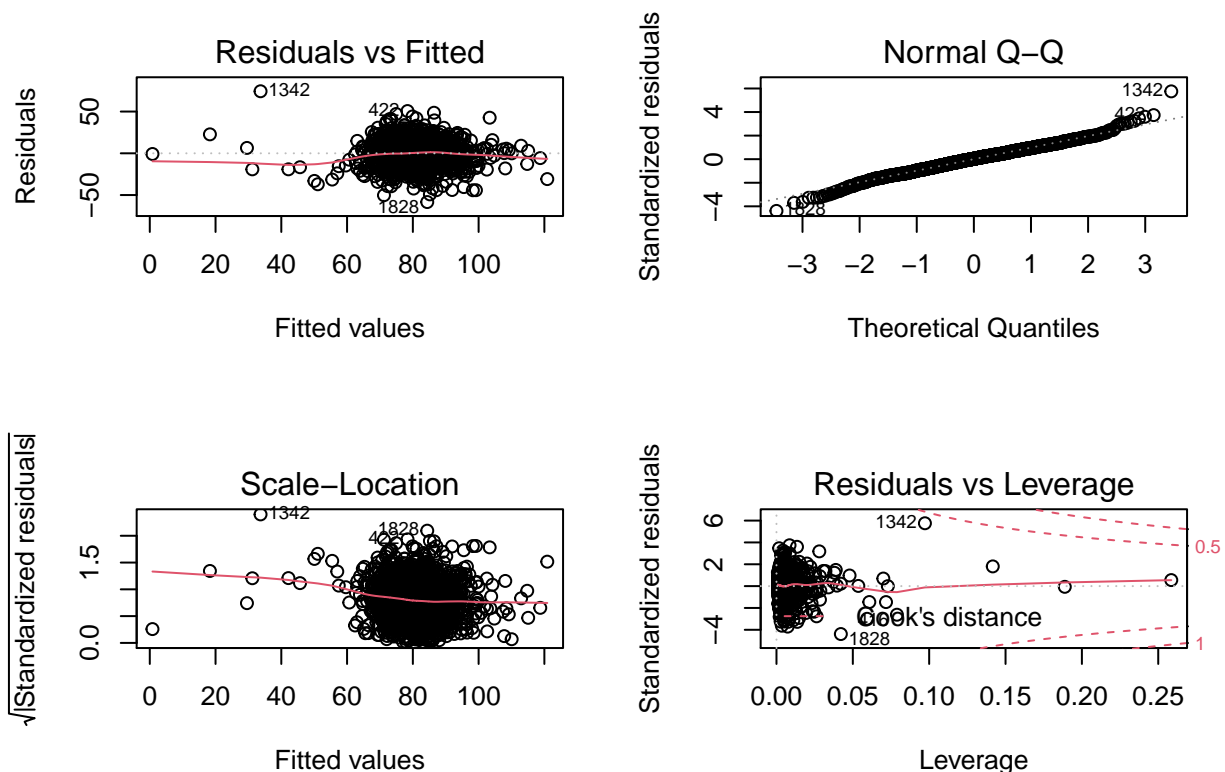
```
##
## Call:
## lm(formula = TARGET_WINS ~ TEAM_BATTING_H + TEAM_BATTING_2B +
##     TEAM_BATTING_3B + TEAM_BATTING_BB + TEAM_BATTING_SO, data = trainData)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -63.489  -8.737   0.549   8.865  51.831
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -16.507104   4.906108  -3.365 0.000782 ***
## TEAM_BATTING_H    0.050124   0.003724  13.460 < 2e-16 ***
```

```
## TEAM_BATTING_2B -0.018508  0.010346 -1.789 0.073804 .
## TEAM_BATTING_3B  0.075932  0.016238  4.676 3.14e-06 ***
## TEAM_BATTING_BB  0.031327  0.002939 10.660 < 2e-16 ***
## TEAM_BATTING_SO  0.011561  0.002065  5.597 2.51e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.79 on 1818 degrees of freedom
## Multiple R-squared:  0.2317, Adjusted R-squared:  0.2296
## F-statistic: 109.7 on 5 and 1818 DF, p-value: < 2.2e-16
```

Model 3: Adding in a Few More

At this point, we're looking for something better than our first model. Steadily, we're adding additional features to see if we can break and R^2 of 30%. Adding in these additional features increased the value from our previous model, but not better than the kitchen sink.

$$R^2 = 0.230$$



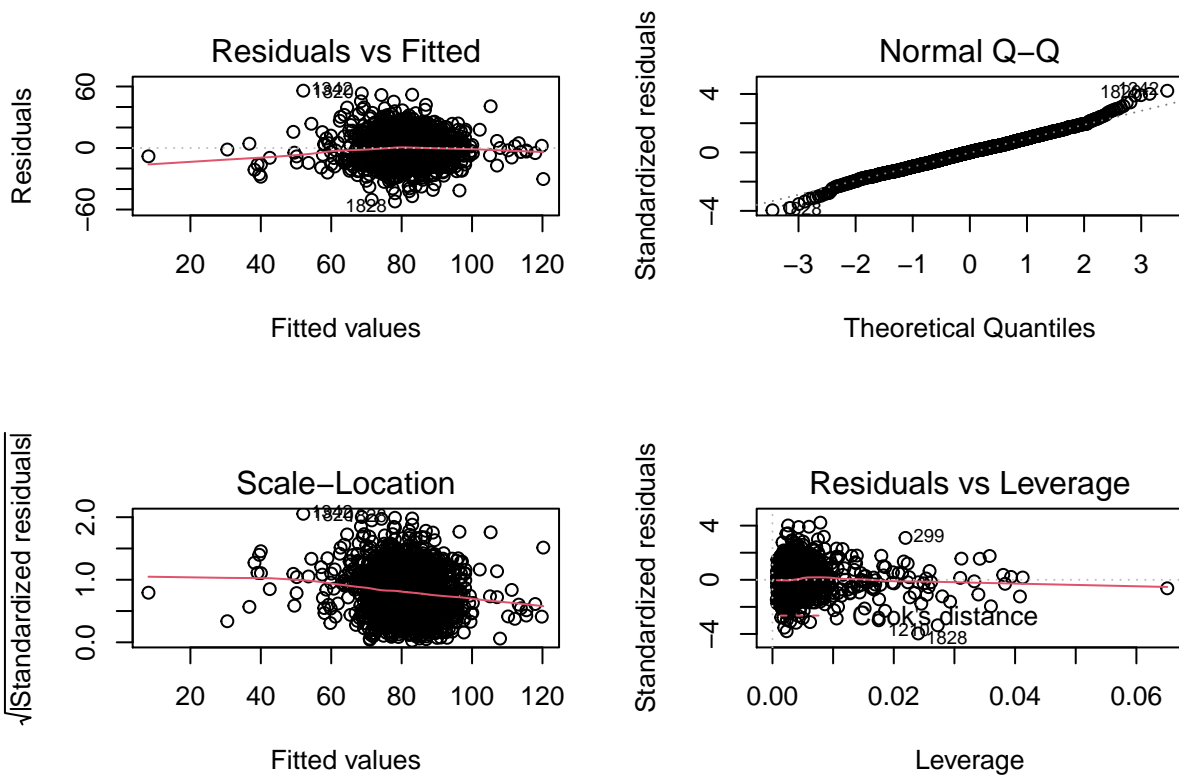
```
##
## Call:
## lm(formula = TARGET_WINS ~ TEAM_BATTING_H + TEAM_BATTING_2B +
##     TEAM_BATTING_3B + TEAM_BATTING_BB + TEAM_BATTING_SO + TEAM_BASERUN_SB +
##     TEAM_BASERUN_CS + TEAM_PITCHING_H + TEAM_PITCHING_HR, data = trainData)
##
```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -58.385  -8.926   0.477   8.994  74.310
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.1535194   5.8554891    0.197  0.84385
## TEAM_BATTING_H    0.0433115   0.0041997   10.313 < 2e-16 ***
## TEAM_BATTING_2B  -0.0084978   0.0104092   -0.816  0.41439
## TEAM_BATTING_3B    0.0573080   0.0190131    3.014  0.00261 **
## TEAM_BATTING_BB    0.0183362   0.0033997    5.393 7.82e-08 ***
## TEAM_BATTING_SO    0.0017697   0.0025771    0.687  0.49237
## TEAM_BASERUN_SB    0.0204570   0.0046655    4.385 1.23e-05 ***
## TEAM_BASERUN_CS    0.0094712   0.0178925    0.529  0.59664
## TEAM_PITCHING_H  -0.0015705   0.0002668  -5.886 4.70e-09 ***
## TEAM_PITCHING_HR   0.0421675   0.0097678    4.317 1.67e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.59 on 1814 degrees of freedom
## Multiple R-squared:  0.255, Adjusted R-squared:  0.2514
## F-statistic: 69.01 on 9 and 1814 DF, p-value: < 2.2e-16
```

Mode 4: Lets start transforming some of these variables

For this model, we attempted to do a few transformations. For three of the features which we found to be left-skewed (TEAM_FIELDING_E, TEAM_BASERUN_SB, and TEAM_BASERUN_CS) we add in the median value. The R^2 value is a bit better than our earlier attempts, but still not as high as the kitchen sink.

$$R^2 = 0.282$$



```
##
## Call:
## lm(formula = TARGET_WINS ~ TEAM_BATTING_H + TEAM_PITCHING_HR +
##     TEAM_FIELDING_DP + fieldingE + baserunSB + baserunCS, data = trainData_model4)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -52.101  -8.869   0.216   8.282  55.937
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    21.796978    4.002272   5.446 5.85e-08 ***
## TEAM_BATTING_H     0.050854    0.002353  21.609 < 2e-16 ***
## TEAM_PITCHING_HR   0.029207    0.006657   4.387 1.21e-05 ***
## TEAM_FIELDING_DP  -0.101189    0.014093  -7.180 1.01e-12 ***
## fieldingE        -0.022605    0.001754 -12.886 < 2e-16 ***
## baserunSB         0.035437    0.004109   8.625 < 2e-16 ***
## baserunCS        -0.025096    0.017767  -1.412  0.158
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.31 on 1817 degrees of freedom
## Multiple R-squared:  0.2842, Adjusted R-squared:  0.2818
## F-statistic: 120.2 on 6 and 1817 DF, p-value: < 2.2e-16
```

Model Selection

Based on the R^2 value of the models, our kitchen sink model (model 1) seemed to perform best, but by no means perfect for this exercise. With some additional features, perhaps team, and season year we could improve our predictions.

According to our F-statistics, all our models had some effect and the residuals appeared more random throughout, even with some tuning, we couldn't get the R^2 to be higher than our first model. We'll use that for our predictions and conclude the assignment by writing those to a csv.

Code Appendix

Libraries

```
r, warning = FALSE, message = FALSE, echo=FALSE
library(kableExtra) library(tidyverse) library(tidymodels) library(VIM) library(naniar) library(GGally) library(caret) library(psych)
```

Data Import

```
{r echo=FALSE}
urlTraining = "https://raw.githubusercontent.com/MsQCompSci/Data621Group4/main/HW1/moneyball-training-data.csv"
rawData <- read.csv(urlTraining)
str(rawData)
```

Train and Test Data Split

```
set.seed(123) trainRowNumbers <- createDataPartition(rawData$INDEX, p=0.8, list=FALSE) trainData <- rawData[trainRowNumbers,] testData <- rawData[-trainRowNumbers,]
```

Summary Statistics

```
summary(rawData)
describe(rawData)
```

Histograms

```
par(mfrow = c(3,3)) for(i in 2:ncol(rawData)) {#distribution of each variable hist(rawData[[i]], main = colnames(rawData[i]), col = "skyblue")}
```


Box Plots

```
trainDataLONG <- rawData %>% select(-INDEX, -TARGET_WINS) %>% gather(key = Variable, value = Value)
```

```
ggplot(trainDataLONG, aes(Variable, Value, fill = Variable)) + geom_boxplot(outlier.colour="blue", outlier.shape=4, outlier.size=2, show.legend=FALSE) + ylim(0,3000) + theme(axis.text.x = element_text(angle = 90, hjust = 1)) + coord_flip()+ labs(title="Moneyball Data Variables (ylim = 3000)")
```

Data Pre-Processing

Missing Data

```
gg_miss_var(rawData) + labs(title="Moneyball Missing Data")
```

```
data.frame(TEAM_BATTING_HBP = sum(is.na(trainDataTEAM_BATTING_HBP))/nrow(trainData), TEAM_BASERUN_CS = sum(is.na(trainDataTEAM_BASERUN_CS))/nrow(trainData), TEAM_FIELDING_DP = sum(is.na(trainDataTEAM_FIELDING_DP))/nrow(trainData), TEAM_BASERUN_SB = sum(is.na(trainDataTEAM_BASERUN_SB))/nrow(trainData), TEAM_BATTING_SO = sum(is.na(trainDataTEAM_BATTING_SO))/nrow(trainData))
```

Correlation Matrix and Multicollinearity

```
corMat <- cor(trainData, use = "complete.obs")
```

```
ggcorr(trainData)
```

Drop Variables

```
findCorrelation(cor(trainData), cutoff = 0.75, verbose = TRUE, names = TRUE)
```

Imputing Missing Data

```
trainData <- trainData %>% select(-TEAM_BATTING_HR, -TEAM_BATTING_HBP)
```

```
Impute Missing Data based on distribution trainData <- trainData %>% mutate(TEAM_BATTING_SO = ifelse(is.na(trainDataTEAM_BATTING_SO), mean(trainDataTEAM_BATTING_SO, na.rm=TRUE), trainDataTEAM_BATTING_SO), TEAM_PITCHING_SO = ifelse(is.na(trainDataTEAM_PITCHING_SO), mean(trainDataTEAM_PITCHING_SO, na.rm = TRUE), trainDataTEAM_PITCHING_SO), TEAM_FIELDING_DP = ifelse(is.na(trainDataTEAM_FIELDING_DP), median(trainDataTEAM_FIELDING_DP, na.rm=TRUE), trainDataTEAM_FIELDING_DP), TEAM_BASERUN_SB = ifelse(is.na(trainDataTEAM_BASERUN_SB), median(trainDataTEAM_BASERUN_SB, na.rm = TRUE), trainDataTEAM_BASERUN_SB), TEAM_BASERUN_CS = ifelse(is.na(trainData$TEAM_BASERUN_CS),
```

Feature Plots

```
featurePlot(y = unlist(trainData$TARGET_WINS), x = trainData, plot = "scatter", type = c("p", "smooth"), span = .5, layout = c(4, 3))
```

Build Models

Model 1: “The Kitchen Sink”

```
model_1 <- lm(TARGET_WINS ~ ., data = trainData) par(mfrow = c(2,2)) plot(model_1) summary(model_1)
```

Model 2: Targeting Most Impactful Features

```
model_2 <- lm(TARGET_WINS ~ TEAM_BATTING_H + TEAM_BATTING_2B + TEAM_BATTING_3B + TEAM_BATTING_BB + TEAM_BATTING_SO, data = trainData) par(mfrow = c(2,2)) plot(model_2) summary(model_2)
```

Model 3: Adding in a Few More

```
model_3 <- lm(TARGET_WINS ~ TEAM_BATTING_H + TEAM_BATTING_2B + TEAM_BATTING_3B + TEAM_BATTING_BB + TEAM_BATTING_SO + TEAM_BASERUN_SB + TEAM_BASERUN_CS + TEAM_PITCHING_H + TEAM_PITCHING_HR, data = trainData) par(mfrow = c(2,2)) plot(model_3) summary(model_3)
```

Model 4: Lets start transforming some of these variables

```
fieldingE <- trainData$TEAM_FIELDING_E + median(trainData$TEAM_FIELDING_E)
baserunSB <- trainData$TEAM_BASERUN_SB + median(trainData$TEAM_BASERUN_SB)
baserunCS <- trainData$TEAM_BASERUN_CS + median(trainData$TEAM_BASERUN_CS)

trainData_model4 <- mutate(trainData, fieldingE = unlist(fieldingE), TEAM_BASERUN_SB = baserunSB, TEAM_BASERUN_CS = baserunCS)

model_4 <- lm(TARGET_WINS ~ TEAM_BATTING_H + TEAM_PITCHING_HR + TEAM_FIELDING_DP + fieldingE + baserunSB + baserunCS, data = trainData_model4) par(mfrow = c(2,2)) plot(model_4) summary(model_4)
```

Model Selection and Write Predictions