

# Data 621 Homework 3

Layla Quinones

10/24/2021

## Libraries

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.0.5
```

```
## Warning: package 'ggplot2' was built under R version 4.0.5
```

```
## Warning: package 'tibble' was built under R version 4.0.5
```

```
## Warning: package 'tidyr' was built under R version 4.0.5
```

```
## Warning: package 'readr' was built under R version 4.0.5
```

```
## Warning: package 'dplyr' was built under R version 4.0.5
```

```
## Warning: package 'forcats' was built under R version 4.0.5
```

```
library(ggplot2)  
library(VIM)
```

```
## Warning: package 'VIM' was built under R version 4.0.5
```

```
## Warning: package 'colorspace' was built under R version 4.0.5
```

```
library(GGally)
```

```
## Warning: package 'GGally' was built under R version 4.0.5
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 4.0.5
```

```
library(broom)
```

## EDA

```
# Load data
# Training
rawTrain <- read.csv("https://raw.githubusercontent.com/MsQCompSci/Data621Group4/main/HW3/crime-training.csv")

#Testing data
rawTest <- read.csv("https://raw.githubusercontent.com/MsQCompSci/Data621Group4/main/HW3/crime-evaluation.csv")

# check to see if we need to clean the data
# gives us a sense of what each predictor is
glimpse(rawTrain)
```

```
## Rows: 466
## Columns: 13
## $ zn      <dbl> 0, 0, 0, 30, 0, 0, 0, 0, 0, 80, 22, 0, 0, 22, 0, 0, 100, 20, 0~
## $ indus   <dbl> 19.58, 19.58, 18.10, 4.93, 2.46, 8.56, 18.10, 18.10, 5.19, 3.6~
## $ chas    <int> 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
## $ nox     <dbl> 0.605, 0.871, 0.740, 0.428, 0.488, 0.520, 0.693, 0.693, 0.515, ~
## $ rm      <dbl> 7.929, 5.403, 6.485, 6.393, 7.155, 6.781, 5.453, 4.519, 6.316, ~
## $ age     <dbl> 96.2, 100.0, 100.0, 7.8, 92.2, 71.3, 100.0, 100.0, 38.1, 19.1, ~
## $ dis     <dbl> 2.0459, 1.3216, 1.9784, 7.0355, 2.7006, 2.8561, 1.4896, 1.6582, ~
## $ rad     <int> 5, 5, 24, 6, 3, 5, 24, 24, 5, 1, 7, 5, 24, 7, 3, 3, 5, 5, 24, ~
## $ tax     <int> 403, 403, 666, 300, 193, 384, 666, 666, 224, 315, 330, 398, 66~
## $ ptratio <dbl> 14.7, 14.7, 20.2, 16.6, 17.8, 20.9, 20.2, 20.2, 20.2, 16.4, 19~
## $ lstat   <dbl> 3.70, 26.82, 18.85, 5.19, 4.82, 7.67, 30.59, 36.98, 5.68, 9.25~
## $ medv    <dbl> 50.0, 13.4, 15.4, 23.7, 37.9, 26.5, 5.0, 7.0, 22.2, 20.9, 24.8~
## $ target  <int> 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, ~
```



```
# All variables are numeric
# categorical variables
# chas

#dicrete
#rad, zn, tax

#all others are continuous
```

## No Missing Values


```
#plot missing values using VIM package
aggr(rawTrain , col=c('navyblue','red'), numbers=TRUE, sortVars=TRUE, labels=names(rawTrain), cex.axis=
```



```
##
## Variables sorted by number of missings:
## Variable Count
##      zn      0
##     indus    0
##      chas    0
##      nox     0
##       rm     0
##      age     0
##      dis     0
##      rad     0
##      tax     0
##   ptratio    0
##     lstat    0
##     medv     0
##    target    0
```

## Correlation

```
#correlation matrix for predictors
ggcorr(rawTrain%>% select(zn:medv))
```



```
#Lets look at some highly correlated variables and drop them
findCorrelation(cor(rowTrain%>% select(zn:medv)),
               cutoff = 0.75,
               verbose = TRUE,
               names = TRUE)
```

```
## Compare row 2 and column 4 with corr 0.76
## Means: 0.539 vs 0.416 so flagging column 2
## Compare row 4 and column 7 with corr 0.769
## Means: 0.487 vs 0.395 so flagging column 4
## Compare row 9 and column 8 with corr 0.906
## Means: 0.46 vs 0.377 so flagging column 9
## Compare row 6 and column 7 with corr 0.751
## Means: 0.417 vs 0.357 so flagging column 6
## All correlations <= 0.75
```

```
## [1] "indus" "nox" "tax" "age"
```

```
# There are 4 highly correlated variables
# I will drop the highest one which is tax which seems to be the most highly correlated
#tax and rad are 0.9 correlated lets look at their relationship to the predictor to see which one to drop
```


## Distribution of Predictors

ADD VARIANCE AND INFLATION FACTORS TO THIS SECTION?

```

par(mfrow = c(3,3))
for(i in 1:ncol(rawTrain)) {#distribution of each variable
  hist(rawTrain[[i]], main = colnames(rawTrain[i]), col = "red")
}

```




```

#binomial data
# indus, tax and rad

#all other variables are skewed except RM


```



## Box Plots

```
#make long
#tax and age has a much different scale so we are seperating it here
longData <- rawTrain %>%
  select(-tax, -age) %>%
  gather(key = Variable, value = Value)

# generate boxplot to identify outliers
ggplot(longData, aes(Variable, Value, fill = Variable)) +
  geom_boxplot(outlier.colour="blue",
               outlier.shape=4,
               outlier.size=2,
               show.legend=FALSE) +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  coord_flip()+
  labs(title="Crime Data Variables", y="Value")
```




*#we can see that zn, medv and lstat has MANY outliers*

```
#make long
#tax and age has a much different scale so we are seperating it here
longData <- rawTrain %>%
  select(tax, age) %>%
  gather(key = Variable, value = Value)

# generate boxplot to identify outliers
ggplot(longData, aes(Variable, Value, fill = Variable)) +
  geom_boxplot(outlier.colour="blue",
               outlier.shape=4,
               outlier.size=2,
               show.legend=FALSE) +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  coord_flip()+
  labs(title="Crime Data Variables", y="Value")
```

## Crime Data Variables



```
# no outliers for tax and age
```

```
#Train/Test Split
```

```
dt <- createDataPartition(iris$Species, p = .8,  
                           list = FALSE,  
                           times = 1)  
train<-rawTrain[dt,]  
test<-rawTrain[-dt,]
```

## Model Building

```
#remove Tax due to high correlation with other variables  
modelOne <- glm(target ~ zn + indus + chas + nox + rm + age + dis + rad + ptratio + lstat + medv , data = train)  
modelOne
```

```
##  
## Call: glm(formula = target ~ zn + indus + chas + nox + rm + age + dis +  
##       rad + ptratio + lstat + medv, family = "binomial", data = train)  
##  
## Coefficients:
```



```
## (Intercept)          zn          indus          chas          nox          rm
##   -49.98632      -0.06821      0.01524      1.68171      62.41402     -2.12008
##          age          dis          rad          ptratio          lstat          medv
##    0.03982      1.05296      0.74014      0.57939      -0.02528      0.36244
##
## Degrees of Freedom: 119 Total (i.e. Null); 108 Residual
## Null Deviance:      163.6
## Residual Deviance: 50.73      AIC: 74.73
```

```
#remove Tax squared age and log lstat
```

```
modelTwo <- glm(target ~ zn + indus + chas + nox + rm + age^2 + dis + rad + ptratio + log2(lstat) + medv, data = train)
modelTwo
```

```
##
## Call:  glm(formula = target ~ zn + indus + chas + nox + rm + age^2 +
##        dis + rad + ptratio + log2(lstat) + medv, family = "binomial",
##        data = train)
##
## Coefficients:
## (Intercept)          zn          indus          chas          nox          rm
##   -45.70491      -0.07452      0.02560      1.71374      63.05722     -2.59852
##          age          dis          rad          ptratio  log2(lstat)          medv
##    0.04930      1.10615      0.75558      0.57837      -0.86561      0.36131
##
## Degrees of Freedom: 119 Total (i.e. Null); 108 Residual
## Null Deviance:      163.6
## Residual Deviance: 50.29      AIC: 74.29
```

```
#This one has a litter lower AIC
```

```
summary(modelTwo)
```

```
##
## Call:
## glm(formula = target ~ zn + indus + chas + nox + rm + age^2 +
##      dis + rad + ptratio + log2(lstat) + medv, family = "binomial",
##      data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.98060  -0.24172  -0.01038   0.00038   2.51026
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -45.70491    19.85879  -2.301  0.02136 *
## zn          -0.07452     0.06833  -1.091  0.27548
## indus         0.02560     0.11069   0.231  0.81708
## chas         1.71374     1.51011   1.135  0.25644
## nox         63.05722    21.23483   2.970  0.00298 **
## rm          -2.59852     1.79138  -1.451  0.14690
## age          0.04930     0.03390   1.454  0.14584
## dis          1.10615     0.55339   1.999  0.04562 *
```

```
## rad          0.75558    0.29663    2.547    0.01086 *
## ptratio      0.57837    0.29231    1.979    0.04786 *
## log2(lstat) -0.86561    1.24316   -0.696    0.48624
## medv         0.36131    0.15789    2.288    0.02212 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 163.645 on 119 degrees of freedom
## Residual deviance: 50.286 on 108 degrees of freedom
## AIC: 74.286
##
## Number of Fisher Scoring iterations: 9
```

```
#log10(zn + 1), log10(dis) and deleted log2(lstat) - not significant
```

```
modelThree <- glm(target ~ log10(zn + 1) + indus + chas + nox + rm + age^2 + log10(dis) + rad + ptratio
```

```
modelThree
```

```
##
## Call:  glm(formula = target ~ log10(zn + 1) + indus + chas + nox + rm +
##       age^2 + log10(dis) + rad + ptratio + medv, family = "binomial",
##       data = train)
##
## Coefficients:
## (Intercept) log10(zn + 1)          indus          chas          nox
## -54.26269    -0.54972         0.05221         1.41733        63.80368
##          rm          age    log10(dis)          rad          ptratio
## -1.92656     0.03476         9.71029         0.77699         0.60580
##          medv
## 0.36206
##
## Degrees of Freedom: 119 Total (i.e. Null); 109 Residual
## Null Deviance: 163.6
## Residual Deviance: 50.47 AIC: 72.47
```

```
#AIC is lower again (not sure if age^2 is helpful)
```

```
summary(modelThree)
```

```
##
## Call:
## glm(formula = target ~ log10(zn + 1) + indus + chas + nox + rm +
##       age^2 + log10(dis) + rad + ptratio + medv, family = "binomial",
##       data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.93941  -0.24838  -0.02038   0.00033   2.59167
##
## Coefficients:
```

```
##           Estimate Std. Error z value Pr(>|z|)
## (Intercept) -54.26269   19.75958  -2.746  0.00603 **
## log10(zn + 1) -0.54972    0.96504  -0.570  0.56892
## indus       0.05221    0.10752   0.486  0.62725
## chas        1.41733    1.50486   0.942  0.34627
## nox        63.80368   21.05195   3.031  0.00244 **
## rm        -1.92656    1.51643  -1.270  0.20392
## age         0.03476    0.02594   1.340  0.18025
## log10(dis)   9.71029    4.81949   2.015  0.04393 *
## rad         0.77699    0.31210   2.490  0.01279 *
## ptratio     0.60580    0.30629   1.978  0.04794 *
## medv        0.36206    0.15616   2.318  0.02042 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 163.645  on 119  degrees of freedom
## Residual deviance:  50.466  on 109  degrees of freedom
## AIC: 72.466
##
## Number of Fisher Scoring iterations: 9
```

*#combine rad and rm (multiplied) - they seemed to correspond in their distributions*

```
modelFour<- glm(target ~ log10(zn + 1) + indus + chas + nox + age^2 + log10(dis) + rad*rm + ptratio +
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
modelFour
```

```
##
## Call:  glm(formula = target ~ log10(zn + 1) + indus + chas + nox + age^2 +
##        log10(dis) + rad * rm + ptratio + medv, family = "binomial",
##        data = train)
##
## Coefficients:
## (Intercept)  log10(zn + 1)          indus          chas          nox
##    -39.69026    -0.39952         0.05430         0.92790        80.55149
##          age    log10(dis)          rad          rm          ptratio
##     0.05098     11.14932     -3.50163     -7.99103         0.90573
##          medv          rad:rm
##     0.60249         0.72035
##
## Degrees of Freedom: 119 Total (i.e. Null);  108 Residual
## Null Deviance:          163.6
## Residual Deviance: 40.73    AIC: 64.73
```

*#AIC is lower #Not sure what the rationale is for this working but it lowered the AIC number and Resid*

```
summary(modelFour)
```

```
##
```

```
## Call:
## glm(formula = target ~ log10(zn + 1) + indus + chas + nox + age^2 +
##      log10(dis) + rad * rm + ptratio + medv, family = "binomial",
##      data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.24476  -0.14203  -0.00414   0.00482   2.05517
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -39.69026   19.76088  -2.009  0.04459 *
## log10(zn + 1) -0.39952    1.23407  -0.324  0.74613
## indus         0.05430    0.12491   0.435  0.66375
## chas          0.92790    1.63062   0.569  0.56933
## nox          80.55149   25.80599   3.121  0.00180 **
## age           0.05098    0.02896   1.760  0.07834 .
## log10(dis)    11.14932    5.21051   2.140  0.03237 *
## rad          -3.50163    1.26718  -2.763  0.00572 **
## rm           -7.99103    3.14895  -2.538  0.01116 *
## ptratio       0.90573    0.37792   2.397  0.01655 *
## medv          0.60249    0.22238   2.709  0.00674 **
## rad:rm         0.72035    0.25034   2.877  0.00401 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 163.645  on 119  degrees of freedom
## Residual deviance:  40.734  on 108  degrees of freedom
## AIC: 64.734
##
## Number of Fisher Scoring iterations: 9
```

```
#delete indus
```

```
modelFive<-glm(target ~ log10(zn+1)+ nox + age^2 + log10(dis) + rad*rm + ptratio + medv, data = train
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
modelFive
```

```
##
## Call:  glm(formula = target ~ log10(zn + 1) + nox + age^2 + log10(dis) +
##      rad * rm + ptratio + medv, family = "binomial", data = train)
##
## Coefficients:
##      (Intercept)  log10(zn + 1)          nox          age      log10(dis)
##      -36.53958      -0.47587      81.22921      0.05416      10.66210
##           rad           rm      ptratio           medv      rad:rm
##      -3.65291      -8.36956      0.85851      0.61188      0.75086
##
## Degrees of Freedom: 119 Total (i.e. Null);  110 Residual
## Null Deviance:      163.6
## Residual Deviance: 41.49      AIC: 61.49
```

```
#AIC is higher #resiudal deviance is lower
# I looked at the histograms and looked for complementary shapes to decide what to multiply
```

## Variable importance

```
summary(modelFive)
```

```
##
## Call:
## glm(formula = target ~ log10(zn + 1) + nox + age^2 + log10(dis) +
##      rad * rm + ptratio + medv, family = "binomial", data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.37810  -0.16278  -0.00419   0.00439   1.97419
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -36.53958   16.42367  -2.225  0.02609 *
## log10(zn + 1)  -0.47587    1.25809  -0.378  0.70524
## nox           81.22921   24.73815   3.284  0.00103 **
## age           0.05416    0.02809   1.928  0.05389 .
## log10(dis)    10.66210    4.89514   2.178  0.02940 *
## rad          -3.65291    1.30703  -2.795  0.00519 **
## rm           -8.36956    3.06969  -2.727  0.00640 **
## ptratio       0.85851    0.35989   2.385  0.01706 *
## medv          0.61188    0.21533   2.842  0.00449 **
## rad:rm        0.75086    0.25810   2.909  0.00362 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 163.645  on 119  degrees of freedom
## Residual deviance:  41.495  on 110  degrees of freedom
## AIC: 61.495
##
## Number of Fisher Scoring iterations: 9
```

```
#indus and zn are not important
```

```
#multiply ptratio*nox (remove squared from age)
```

```
modelSix<- glm(target ~ log10(zn + 1) + age + ptratio*nox + log10(dis) + rad*rm + medv, data = train, family = "binomial")
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
modelSix
```

```
##
```

```
## Call: glm(formula = target ~ log10(zn + 1) + age + ptratio * nox +
##       log10(dis) + rad * rm + medv, family = "binomial", data = train)
##
## Coefficients:
## (Intercept)  log10(zn + 1)          age          ptratio          nox
##   -50.42921   -0.48978         0.05554         1.58424       105.27247
##   log10(dis)         rad          rm          medv   ptratio:nox
##    10.68244   -3.59743   -8.29831         0.61232        -1.31047
##      rad:rm
##      0.74144
##
## Degrees of Freedom: 119 Total (i.e. Null);  109 Residual
## Null Deviance:      163.6
## Residual Deviance: 41.42    AIC: 63.42
```

*#AIC is lower*

```
summary(modelSix)
```

```
##
## Call:
## glm(formula = target ~ log10(zn + 1) + age + ptratio * nox +
##       log10(dis) + rad * rm + medv, family = "binomial", data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.34744  -0.15485  -0.00417   0.00479   1.98426
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -50.42921    55.18900  -0.914  0.36085
## log10(zn + 1)  -0.48978     1.29626  -0.378  0.70555
## age           0.05554     0.02869   1.936  0.05291 .
## ptratio       1.58424     2.76305   0.573  0.56640
## nox          105.27247    95.21624   1.106  0.26889
## log10(dis)    10.68244     4.93429   2.165  0.03039 *
## rad          -3.59743     1.32868  -2.708  0.00678 **
## rm           -8.29831     3.11600  -2.663  0.00774 **
## medv          0.61232     0.21790   2.810  0.00495 **
## ptratio:nox   -1.31047     4.92343  -0.266  0.79011
## rad:rm        0.74144     0.26133   2.837  0.00455 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 163.645  on 119  degrees of freedom
## Residual deviance:  41.421  on 109  degrees of freedom
## AIC: 63.421
##
## Number of Fisher Scoring iterations: 9
```

## Test Models

*#Make predictions*

```
predOne = predict(modelOne,test, type = "response")
predTwo = predict(modelTwo,test, type = "response")
predThree = predict(modelThree,test, type = "response")
predFour = predict(modelFour,test, type = "response")
predFive = predict(modelFive,test, type = "response")
predSix = predict(modelSix,test, type = "response")
```

*#Error Measures*

```
data.frame(modelOne = postResample(pred = predOne, obs = test$target), modelTwo = postResample(pred = p
```

```
##           modelOne  modelTwo modelThree modelFour modelFive  modelSix
## RMSE      0.2787735 0.2791040  0.2725896 0.2868524 0.2814246 0.2796029
## Rsquared  0.6950779 0.6948205  0.7076774 0.6806475 0.6914870 0.6952936
## MAE       0.1253909 0.1253186  0.1218843 0.1217098 0.1197176 0.1185218
```

*#We can see RMSE is increasing which means the fit is better for every model - This doesnt reflect very*

## Confusion Matrix and Accuracy Measurment

*#Extract Accuracy*

*#Model One*

*#format predictions to binary*

```
resultsFitOne <- ifelse(predOne > 0.5,1,0)
resultsFitOne <- as.factor(resultsFitOne)
```

*#Confusion Matrix to Extract Accuracy*

```
cOne <- confusionMatrix(as.factor(test$target),resultsFitOne)
accOne <- as.data.frame(cOne$overall)[1]
accOne<- accOne %>%
  slice(1)
```

*#Model Two*

*#format predictions to binary*

```
resultsFitTwo <- ifelse(predTwo > 0.5,1,0)
resultsFitTwo <- as.factor(resultsFitTwo)
```

*#Confusion Matrix to Extract Accuracy*

```
cTwo <- confusionMatrix(resultsFitTwo, as.factor(test$target))
accTwo <- as.data.frame(cTwo$overall)[1]
accTwo<- accTwo %>%
  slice(1)
```

*#Model Three*

*#format predictions to binary*

```

resultsFitThree<- ifelse(predThree > 0.5,1,0)
resultsFitThree <- as.factor(resultsFitThree)

#Confusion Matrix to Extract Accuracy
cThree <- confusionMatrix(resultsFitThree, as.factor(test$target))
accThree <- as.data.frame(cThree$overall)[1]
accThree<- accThree%>%
  slice(1)

#Model Four
#format predictions to binary
resultsFitFour<- ifelse(predFour > 0.5,1,0)
resultsFitFour <- as.factor(resultsFitFour)

#Confusion Matrix to Extract Accuracy
cFour <- confusionMatrix(resultsFitFour, as.factor(test$target))
accFour <- as.data.frame(cFour$overall)[1]
accFour<- accFour%>%
  slice(1)

#Model Five
#format predictions to binary
resultsFitFive<- ifelse(predFive > 0.5,1,0)
resultsFitFive <- as.factor(resultsFitFive)

#Confusion Matrix to Extract Accuracy
cFive <- confusionMatrix(resultsFitFive, as.factor(test$target))
accFive <- as.data.frame(cFive$overall)[1]
accFive<- accFive%>%
  slice(1)

#Model Six
#format predictions to binary
resultsFitSix<- ifelse(predSix > 0.5,1,0)
resultsFitSix <- as.factor(resultsFitSix)

#Confusion Matrix to Extract Accuracy
cSix<- confusionMatrix(resultsFitSix, as.factor(test$target))
accSix <- as.data.frame(cSix$overall)[1]
accSix<- accSix%>%
  slice(1)

#create a table with accuracies
data.frame(c(accOne, accTwo, accThree, accFour, accFive, accSix))

##   cOne.overall cTwo.overall cThree.overall cFour.overall cFive.overall
## 1    0.8843931    0.8815029    0.8959538    0.8901734    0.8988439
##   cSix.overall
## 1    0.9017341

```



*#Here we see that our best models are Five and Six in terms of accuracy*

WE NEED QQ PLOTS OR SOME OTHER VISUAL TO HELP US TALK ABOUT GOODNESS OF FIT GETTING HIGHER ALTHOUGH THE ACCURACY IS NOT CHANGING SO WE CAN CHOOSE ONE (FIVE OR SIX)

AUC or ROC curve