

Data 621 Homework 3

Layla Quinones

10/24/2021

Libraries

```
library(tidyverse)
library(ggplot2)
library(VIM)
library(GGally)
library(caret)
library(broom)
```

EDA

```
# Load data
# Training
rawTrain <- read.csv("https://raw.githubusercontent.com/MsQCompSci/Data621Group4/main/HW3/crime-training")

#Testing data
rawTest <- read.csv("https://raw.githubusercontent.com/MsQCompSci/Data621Group4/main/HW3/crime-evaluation")

# check to see if we need to clean the data
# gives us a sense of what each predictor is
glimpse(rawTrain)
```

```
## Rows: 466
## Columns: 13
## $ zn      <dbl> 0, 0, 0, 30, 0, 0, 0, 0, 0, 80, 22, 0, 0, 22, 0, 0, 100, 20...
## $ indus   <dbl> 19.58, 19.58, 18.10, 4.93, 2.46, 8.56, 18.10, 18.10, 5.19, ...
## $ chas    <int> 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ nox     <dbl> 0.605, 0.871, 0.740, 0.428, 0.488, 0.520, 0.693, 0.693, 0.5...
## $ rm      <dbl> 7.929, 5.403, 6.485, 6.393, 7.155, 6.781, 5.453, 4.519, 6.3...
## $ age     <dbl> 96.2, 100.0, 100.0, 7.8, 92.2, 71.3, 100.0, 100.0, 38.1, 19...
## $ dis     <dbl> 2.0459, 1.3216, 1.9784, 7.0355, 2.7006, 2.8561, 1.4896, 1.6...
## $ rad     <int> 5, 5, 24, 6, 3, 5, 24, 24, 5, 1, 7, 5, 24, 7, 3, 3, 5, 5, 2...
## $ tax     <int> 403, 403, 666, 300, 193, 384, 666, 666, 224, 315, 330, 398, ...
## $ ptratio <dbl> 14.7, 14.7, 20.2, 16.6, 17.8, 20.9, 20.2, 20.2, 20.2, 16.4, ...
## $ lstat   <dbl> 3.70, 26.82, 18.85, 5.19, 4.82, 7.67, 30.59, 36.98, 5.68, 9...
## $ medv    <dbl> 50.0, 13.4, 15.4, 23.7, 37.9, 26.5, 5.0, 7.0, 22.2, 20.9, 2...
## $ target  <int> 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, ...
```

```

# All varaibles are numeric
# categorical variables
# chas

#dicrete
#rad, zn, tax

#all others are continuous


```

No Missing Values

```

#plot missing values using VIM package
aggr(rawTrain , col=c('navyblue','red'), numbers=TRUE, sortVars=TRUE, labels=names(rawTrain), cex.axis=

```



```


##
## Variables sorted by number of missings:
## Variable Count
##      zn      0
##    indus    0
##     chas    0
##     nox     0
##      rm     0
##     age     0

```

```
##      dis      0
##      rad      0
##      tax      0
##  ptratio      0
##      lstat      0
##      medv      0
##      target    0
```

```
#plot missing values using VIM package
```


```
aggr(rawTest , col=c('navyblue','red'), numbers=TRUE, sortVars=TRUE, labels=names(rawTrain), cex.axis=.7)
```




```
##
## Variables sorted by number of missings:
## Variable Count
##      zn      0
##      indus    0
##      chas     0
##      nox      0
##      rm       0
##      age      0
##      dis      0
##      rad      0
##      tax      0
##  ptratio      0
##      lstat    0
##      medv     0
```

Correlation

```
#correlation matrix for predictors  
ggcorr(rawTrain%>% select(zn:medv))
```



```
#Identify highly correlated variables  
ggcorr(rawTrain%>% select(zn:medv))
```



```
#Lets look at some highly correlated variables and drop them
findCorrelation(cor(rawTrain%>% select(zn:medv)),
               cutoff = 0.75,
               verbose = TRUE,
               names = TRUE)
```

```
## Compare row 2 and column 4 with corr 0.76
## Means: 0.539 vs 0.416 so flagging column 2
## Compare row 4 and column 7 with corr 0.769
## Means: 0.487 vs 0.395 so flagging column 4
## Compare row 9 and column 8 with corr 0.906
## Means: 0.46 vs 0.377 so flagging column 9
## Compare row 6 and column 7 with corr 0.751
## Means: 0.417 vs 0.357 so flagging column 6
## All correlations <= 0.75
```

```
## [1] "indus" "nox" "tax" "age"
```

```
# There are 4 highly correlated variables
# I will drop the highest one which is tax which seems to be the most highly correlated
#tax and rad are 0.9 correlated lets look at their relationship to the predictor to see which one to drop
```

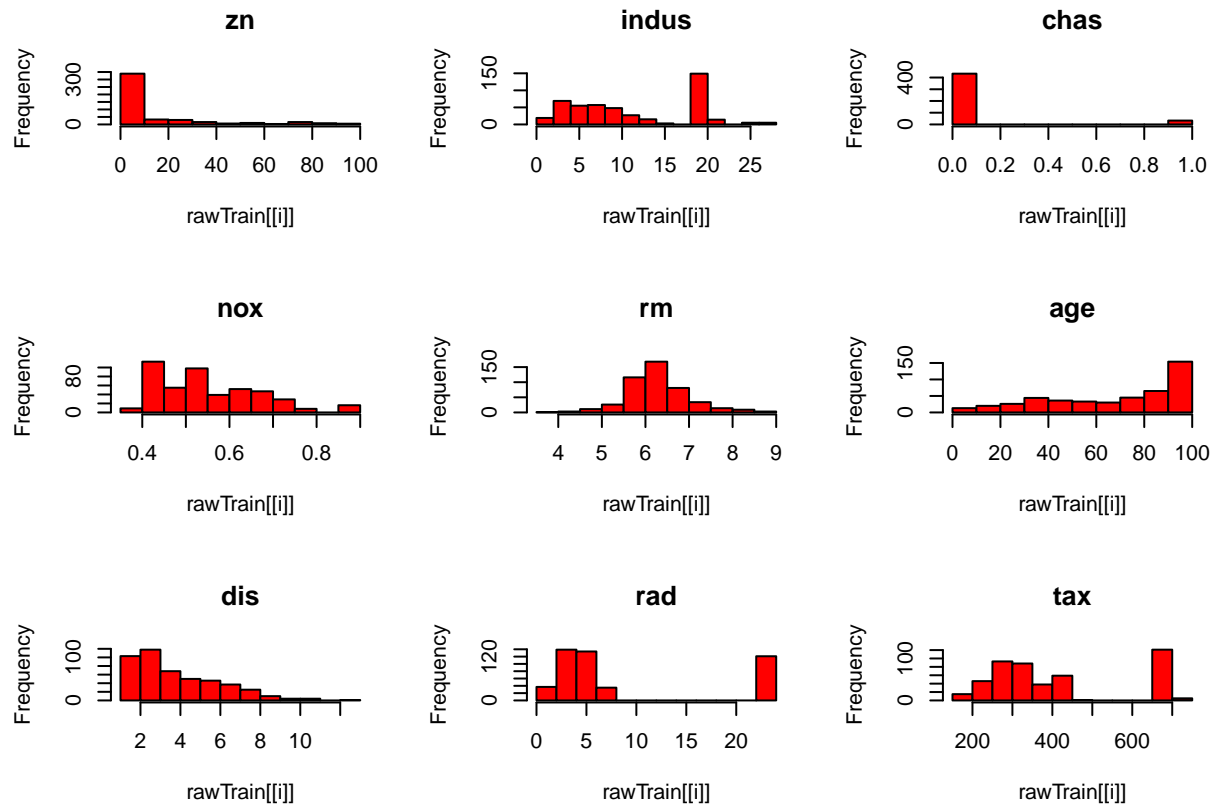
Distribution of Predictors

ADD VARIANCE AND INFLATION FACTORS TO THIS SECTION

```

par(mfrow = c(3,3))
for(i in 1:ncol(rawTrain)) {#distribution of each variable
  hist(rawTrain[[i]], main = colnames(rawTrain[i]), col = "red")
}

```




```

#binomial data
# indus, tax and rad

#all other variables are skewed except RM

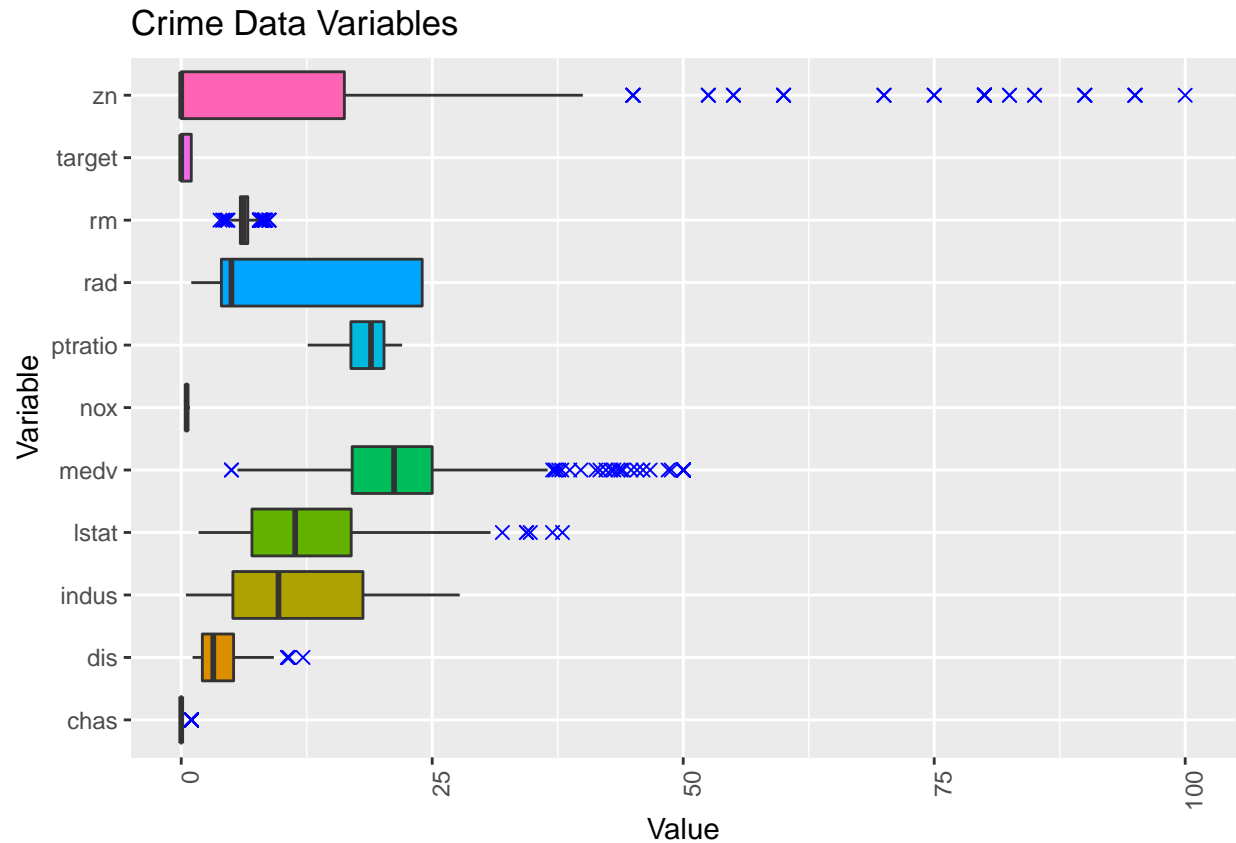
```



Box Plots

```
#make long
#tax and age has a much different scale so we are seperating it here
longData <- rawTrain %>%
  select(-tax, -age) %>%
  gather(key = Variable, value = Value)

# generate boxplot to identify outliers
ggplot(longData, aes(Variable, Value, fill = Variable)) +
  geom_boxplot(outlier.colour="blue",
               outlier.shape=4,
               outlier.size=2,
               show.legend=FALSE) +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  coord_flip()+
  labs(title="Crime Data Variables", y="Value")
```




#we can see that zn, medv and lstat has MANY outliers

```
#make long
#tax and age has a much different scale so we are seperating it here
longData <- rawTrain %>%
  select(tax, age) %>%
  gather(key = Variable, value = Value)

# generate boxplot to identify outliers
ggplot(longData, aes(Variable, Value, fill = Variable)) +
  geom_boxplot(outlier.colour="blue",
               outlier.shape=4,
               outlier.size=2,
               show.legend=FALSE) +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  coord_flip()+
  labs(title="Crime Data Variables", y="Value")
```


Crime Data Variables



```
# no outliers for tax and age
```

```
#Train/Test Split
dt = sort(sample(nrow(rawTrain), nrow(rawTrain)*.8))
train<-rawTrain[dt,]
test<-rawTrain[-dt,]
```

Model Building

```
#remove Tax due to high correlation with other variables
modelOne <- glm(target ~ zn + indus + chas + nox + rm + age + dis + rad + ptratio + lstat + medv , data = train)
modelOne
```

```
##
## Call: glm(formula = target ~ zn + indus + chas + nox + rm + age + dis +
##       rad + ptratio + lstat + medv, family = "binomial", data = train)
##
## Coefficients:
## (Intercept)          zn          indus          chas          nox          rm
## -36.68816    -0.06465   -0.04780     1.18303    40.32164   -0.56808
##          age          dis          rad          ptratio          lstat          medv
```


```
##      0.02676      0.71335      0.45953      0.36870      0.06568      0.19556
##
## Degrees of Freedom: 371 Total (i.e. Null); 360 Residual
## Null Deviance:      515
## Residual Deviance: 170.1      AIC: 194.1
```

```
# squared transformation to age and lstat
```


```
#age before squared
summary(train$age)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      2.90  40.95   76.60   67.61   93.83  100.00
```

```
#age before squared
qplot(train$age)
```




```
#age after squared
qplot((train$age)^(2))
```




```
#lstat before log
summary(train$lstat)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1.730   6.928   11.110   12.586   17.093   37.970
```

```
#lstat before log
qplot(train$lstat)
```



```
#lstat after log  
qplot(log10(train$lstat))
```



```
#remove Tax squared age and log lstat
```

```
modelTwo <- glm(target ~ zn + indus + chas + nox + rm + age2 + dis + rad + ptratio + log10(lstat) + me
```

```
modelTwo
```

```
##
```

```
## Call: glm(formula = target ~ zn + indus + chas + nox + rm + age2 +
```

```
## dis + rad + ptratio + log10(lstat) + medv, family = "binomial",
```

```
## data = train)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)          zn          indus          chas          nox
```

```
## -36.02179    -0.05964    -0.04500     1.27324    40.86438
```

```
##          rm          age          dis          rad          ptratio
```

```
## -0.81350     0.03202     0.73546     0.46057     0.38027
```

```
## log10(lstat)      medv
```

```
##  0.53610      0.19918
```

```
##
```

```
## Degrees of Freedom: 371 Total (i.e. Null);  360 Residual
```

```
## Null Deviance:      515
```

```
## Residual Deviance: 171.3    AIC: 195.3
```

```
#This one has a litter lower AIC
```

```
#remove Tax squared age and log lstat - log dis and zn +1
modelThree <- glm(target ~ log10(zn + 1) + indus + chas + nox + rm + age^2 + log10(dis) + rad + ptratio
modelThree
```

```
##
## Call: glm(formula = target ~ log10(zn + 1) + indus + chas + nox + rm +
##       age^2 + log10(dis) + rad + ptratio + log10(lstat) + medv,
##       family = "binomial", data = train)
##
## Coefficients:
##   (Intercept)  log10(zn + 1)          indus          chas          nox
##   -44.047678    -0.972791    -0.007393     1.079595    46.414029
##           rm          age    log10(dis)          rad          ptratio
##   -0.888001     0.039243     9.603166     0.497630     0.418534
## log10(lstat)          medv
##     0.929415     0.241107
##
## Degrees of Freedom: 371 Total (i.e. Null);  360 Residual
## Null Deviance:      515
## Residual Deviance: 164.5      AIC: 188.5
```

```
#AIC is lower again
```

```
#add lstat*age
modelFour <- glm(target ~ log10(zn+ 1) + indus + nox + rm + log10(dis) + rad + ptratio + medv + lstat
modelFour
```

```
##
## Call: glm(formula = target ~ log10(zn + 1) + indus + nox + rm + log10(dis) +
##       rad + ptratio + medv + lstat * age + age^2 + log10(lstat) +
##       chas, family = "binomial", data = rawTrain)
##
## Coefficients:
##   (Intercept)  log10(zn + 1)          indus          nox          rm
##   -47.867344    -0.971631    -0.079895     57.940008    -0.955576
##   log10(dis)          rad          ptratio          medv          lstat
##   10.727914     0.600422     0.455900     0.211453     0.533072
##           age    log10(lstat)          chas    lstat:age
##     0.082900    -7.758868     1.178626    -0.003236
##
## Degrees of Freedom: 465 Total (i.e. Null);  452 Residual
## Null Deviance:      645.9
## Residual Deviance: 181.2      AIC: 209.2
```

```
#Here I decided to take lstat and age and multiply them because age is highly correlated and lstat is s
```

Test Models

```
#Make predictions
predOne = predict(modelOne,test, type = "response")
predTwo = predict(modelTwo,test, type = "response")
predThree = predict(modelThree,test, type = "response")
predFour = predict(modelFour,test, type = "response")
```

```
#measure accuracy
postResample(pred = predOne, obs = test$target)
```

```
##      RMSE  Rsquared      MAE
## 0.2249477 0.7997835 0.1203042
```

```
#measure accuracy
postResample(pred = predTwo, obs = test$target)
```

```
##      RMSE  Rsquared      MAE
## 0.2170584 0.8144284 0.1160138
```

```
#measure accuracy
postResample(pred = predThree, obs = test$target)
```

```
##      RMSE  Rsquared      MAE
## 0.2141277 0.8189267 0.1118686
```

```
#measure accuracy
postResample(pred = predFour, obs = test$target)
```

```
##      RMSE  Rsquared      MAE
## 0.17975252 0.87386770 0.09000409
```

Confusion Matric and Accuracy Measurment

```
resultsFit<- ifelse(predOne > 0.5,1,0)
resultsFit <- as.factor(resultsFit)
#confusionMatrix(test$target, resultsFit)
resultsFit
```

```
##      1   4   7  18  25  43  47  49  66  67  77  79  81  92  94 104 109 111 116 117
##      1   0   1   1   0   0   1   1   1   0   1   0   0   0   0   0   1   1   0   1
## 119 122 124 130 146 151 153 154 161 167 181 182 183 184 186 190 195 200 221 227
##      1   0   1   0   1   1   0   1   1   0   0   0   1   0   1   1   1   1   1   1
## 228 229 231 232 243 244 257 260 262 263 265 267 271 276 284 286 295 298 301 303
##      1   1   1   1   1   1   0   1   0   1   1   1   1   1   1   1   1   1   0   0
## 312 314 324 327 331 348 349 351 352 361 363 365 366 372 379 387 391 393 396 400
##      1   0   0   0   1   0   0   0   0   0   0   1   0   0   1   1   0   1   0   1
## 405 410 411 416 419 425 426 437 441 442 454 455 460 464
##      1   0   0   0   1   0   0   0   0   1   1   1   0   1
## Levels: 0 1
```

Anova Tests for each model

```
#Looking at strength of variables  
anova(modelOne, test = 'Chisq')
```

```
## Analysis of Deviance Table  
##  
## Model: binomial, link: logit  
##  
## Response: target  
##  
## Terms added sequentially (first to last)  
##  
##  
##          Df Deviance Resid. Df Resid. Dev  Pr(>Chi)  
## NULL                                371    515.01  
## zn          1  101.237      370    413.78 < 2.2e-16 ***  
## indus       1   85.715      369    328.06 < 2.2e-16 ***  
## chas        1    2.602      368    325.46  0.106754  
## nox         1   99.035      367    226.42 < 2.2e-16 ***  
## rm          1    1.391      366    225.03  0.238154  
## age         1    0.080      365    224.95  0.777458  
## dis         1    5.701      364    219.25  0.016957 *  
## rad         1   34.639      363    184.61 3.969e-09 ***  
## ptratio    1    3.524      362    181.09  0.060475 .  
## lstat      1    1.459      361    179.63  0.227119  
## medv       1    9.503      360    170.13  0.002051 **  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#Looking at strength of variables  
anova(modelTwo, test = 'Chisq')
```

```
## Analysis of Deviance Table  
##  
## Model: binomial, link: logit  
##  
## Response: target  
##  
## Terms added sequentially (first to last)  
##  
##  
##          Df Deviance Resid. Df Resid. Dev  Pr(>Chi)  
## NULL                                371    515.01  
## zn          1  101.237      370    413.78 < 2.2e-16 ***  
## indus       1   85.715      369    328.06 < 2.2e-16 ***  
## chas        1    2.602      368    325.46  0.106754  
## nox         1   99.035      367    226.42 < 2.2e-16 ***  
## rm          1    1.391      366    225.03  0.238154  
## age         1    0.080      365    224.95  0.777458  
## dis         1    5.701      364    219.25  0.016957 *  
## rad         1   34.639      363    184.61 3.969e-09 ***
```



```
## ptratio      1    3.524      362    181.09  0.060475 .
## log10(lstat) 1    0.016      361    181.07  0.898119
## medv         1    9.727      360    171.35  0.001816 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#Looking at strength of variables
anova(modelThree, test = 'Chisq')
```

```
## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: target
##
## Terms added sequentially (first to last)
##
##
##          Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
## NULL                                371     515.01
## log10(zn + 1) 1    93.354      370     421.66 < 2.2e-16 ***
## indus         1    88.185      369     333.47 < 2.2e-16 ***
## chas          1     2.474      368     331.00 0.1157461
## nox           1   104.010      367     226.99 < 2.2e-16 ***
## rm            1     1.301      366     225.69 0.2539696
## age           1     0.136      365     225.55 0.7122263
## log10(dis)    1     8.262      364     217.29 0.0040484 **
## rad           1    36.959      363     180.33 1.206e-09 ***
## ptratio       1     3.041      362     177.29 0.0811920 .
## log10(lstat)  1     0.000      361     177.29 0.9962682
## medv          1    12.827      360     164.46 0.0003417 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

NEXT I WANT TO TRY BOX COX TRANSFORMATIONS on things we deleted?

```
#Looking at strength of variables (now we have all strong variables)
anova(modelFour, test = 'Chisq')
```

```
## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: target
##
## Terms added sequentially (first to last)
##
##
##          Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
## NULL                                465     645.88
## log10(zn + 1) 1   116.753      464     529.12 < 2.2e-16 ***
## indus         1    89.384      463     439.74 < 2.2e-16 ***
## nox           1   155.463      462     284.28 < 2.2e-16 ***
```

```
## rm          1      7.067      461      277.21  0.007852 **
## log10(dis)   1      9.104      460      268.10  0.002550 **
## rad         1     55.030      459      213.07  1.187e-13 ***
## ptratio     1      1.954      458      211.12  0.162162
## medv        1      3.969      457      207.15  0.046359 *
## lstat       1      7.083      456      200.07  0.007781 **
## age         1      9.544      455      190.53  0.002006 **
## log10(lstat) 1      1.693      454      188.83  0.193158
## chas        1      2.154      453      186.68  0.142243
## lstat:age    1      5.446      452      181.23  0.019609 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

WEE NEED QQ PLOTS AND ACCURACY

AUC or ROC curve