

Data 621 Homework 3

Layla Quinones

10/24/2021

Libraries

```
library(tidyverse)
library(ggplot2)
library(VIM)
library(GGally)
library(caret)
```

EDA

```
# Load data
# Training
rawTrain <- read.csv("https://raw.githubusercontent.com/MsQCompSci/Data621Group4/main/HW3/crime-training")

#Testing data
rawTest <- read.csv("https://raw.githubusercontent.com/MsQCompSci/Data621Group4/main/HW3/crime-evaluation")

# check to see if we need to clean the data
# gives us a sense of what each predictor is
glimpse(rawTrain)
```

```
## Rows: 466
## Columns: 13
## $ zn      <dbl> 0, 0, 0, 30, 0, 0, 0, 0, 0, 80, 22, 0, 0, 22, 0, 0, 100, 20...
## $ indus   <dbl> 19.58, 19.58, 18.10, 4.93, 2.46, 8.56, 18.10, 18.10, 5.19, ...
## $ chas    <int> 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ nox     <dbl> 0.605, 0.871, 0.740, 0.428, 0.488, 0.520, 0.693, 0.693, 0.5...
## $ rm      <dbl> 7.929, 5.403, 6.485, 6.393, 7.155, 6.781, 5.453, 4.519, 6.3...
## $ age     <dbl> 96.2, 100.0, 100.0, 7.8, 92.2, 71.3, 100.0, 100.0, 38.1, 19...
## $ dis     <dbl> 2.0459, 1.3216, 1.9784, 7.0355, 2.7006, 2.8561, 1.4896, 1.6...
## $ rad     <int> 5, 5, 24, 6, 3, 5, 24, 24, 5, 1, 7, 5, 24, 7, 3, 3, 5, 5, 2...
## $ tax     <int> 403, 403, 666, 300, 193, 384, 666, 666, 224, 315, 330, 398, ...
## $ ptratio <dbl> 14.7, 14.7, 20.2, 16.6, 17.8, 20.9, 20.2, 20.2, 20.2, 16.4, ...
## $ lstat   <dbl> 3.70, 26.82, 18.85, 5.19, 4.82, 7.67, 30.59, 36.98, 5.68, 9...
## $ medv    <dbl> 50.0, 13.4, 15.4, 23.7, 37.9, 26.5, 5.0, 7.0, 22.2, 20.9, 2...
## $ target  <int> 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, ...
```

```

# All varaibles are numeric
# categorical variables
# chas

#dicrete
#rad, zn, tax

#all others are continuous


```

No Missing Values

```

#plot missing values using VIM package
aggr(rawTrain , col=c('navyblue','red'), numbers=TRUE, sortVars=TRUE, labels=names(rawTrain), cex.axis=

```



```


##
## Variables sorted by number of missings:
## Variable Count
##      zn      0
##    indus    0
##     chas    0
##     nox     0
##      rm     0
##     age     0

```

```
##      dis      0
##      rad      0
##      tax      0
##  ptratio      0
##      lstat      0
##      medv      0
##      target    0
```

```
#plot missing values using VIM package
```


```
aggr(rawTest , col=c('navyblue','red'), numbers=TRUE, sortVars=TRUE, labels=names(rawTrain), cex.axis=.7)
```




```
##
## Variables sorted by number of missings:
## Variable Count
##      zn      0
##      indus    0
##      chas     0
##      nox      0
##      rm       0
##      age      0
##      dis      0
##      rad      0
##      tax      0
##  ptratio      0
##      lstat     0
##      medv      0
```

Correlation

```
#correlation matrix for predictors
ggcorr(rawTrain%>% select(zn:medv))
```



```
#Identify highly correlated variables
ggcorr(rawTrain%>% select(zn:medv))
```



```
#Lets look at some highly correlated variables and drop them
findCorrelation(cor(rawTrain%>% select(zn:medv)),
  cutoff = 0.75,
  verbose = TRUE,
  names = TRUE)
```

```
## Compare row 2 and column 4 with corr 0.76
## Means: 0.539 vs 0.416 so flagging column 2
## Compare row 4 and column 7 with corr 0.769
## Means: 0.487 vs 0.395 so flagging column 4
## Compare row 9 and column 8 with corr 0.906
## Means: 0.46 vs 0.377 so flagging column 9
## Compare row 6 and column 7 with corr 0.751
## Means: 0.417 vs 0.357 so flagging column 6
## All correlations <= 0.75
```

```
## [1] "indus" "nox" "tax" "age"
```

```
# There are 4 highly correlated variables
# I will drop the highest one which is tax which seems to be the most highly correlated
#tax and rad are 0.9 correlated lets look at their relationship to the predictor to see which one to drop
```

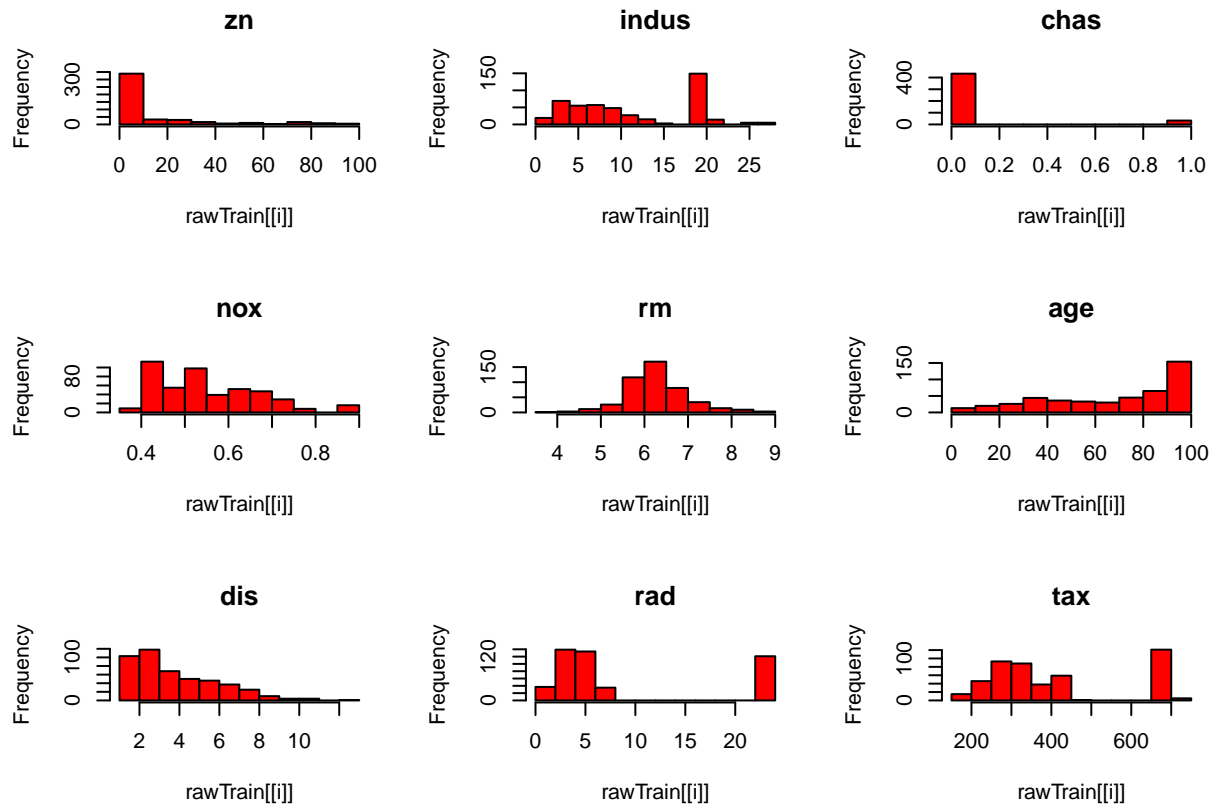
Distribution of Predictors

ADD VARIANCE AND INFLATION FACTORS TO THIS SECTION

```

par(mfrow = c(3,3))
for(i in 1:ncol(rawTrain)) {#distribution of each variable
  hist(rawTrain[[i]], main = colnames(rawTrain[i]), col = "red")
}

```




```

#binomial data
# indus, tax and rad

#all other variables are skewed except RM

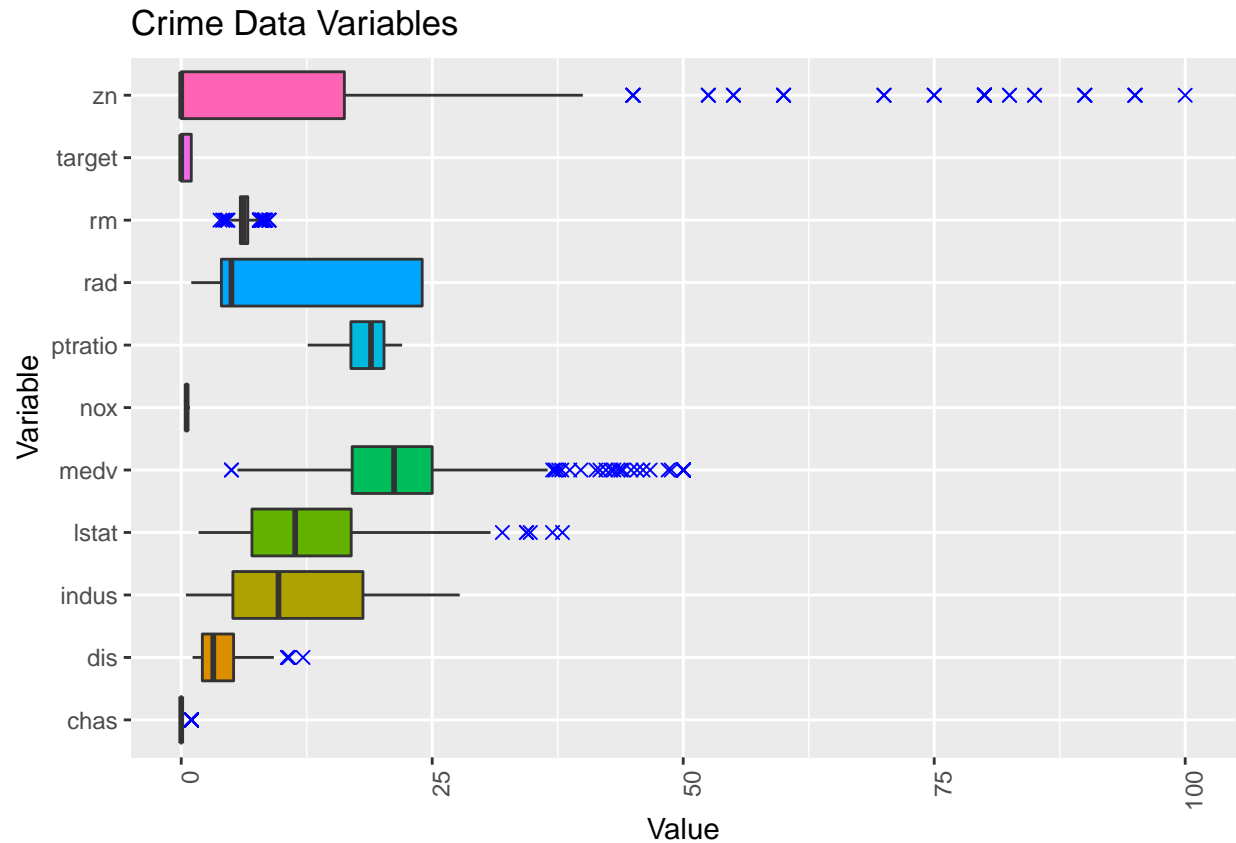
```



Box Plots

```
#make long
#tax and age has a much different scale so we are seperating it here
longData <- rawTrain %>%
  select(-tax, -age) %>%
  gather(key = Variable, value = Value)

# generate boxplot to identify outliers
ggplot(longData, aes(Variable, Value, fill = Variable)) +
  geom_boxplot(outlier.colour="blue",
               outlier.shape=4,
               outlier.size=2,
               show.legend=FALSE) +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  coord_flip()+
  labs(title="Crime Data Variables", y="Value")
```




#we can see that zn, medv and lstat has MANY outliers

```
#make long
#tax and age has a much different scale so we are seperating it here
longData <- rawTrain %>%
  select(tax, age) %>%
  gather(key = Variable, value = Value)

# generate boxplot to identify outliers
ggplot(longData, aes(Variable, Value, fill = Variable)) +
  geom_boxplot(outlier.colour="blue",
               outlier.shape=4,
               outlier.size=2,
               show.legend=FALSE) +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  coord_flip()+
  labs(title="Crime Data Variables", y="Value")
```


Crime Data Variables



```
# no outliers for tax and age
```

Model Building

```
#remove Tax due to high correlation with other variables
model0ne <- glm(target ~ zn + indus + chas + nox + rm + age + dis + rad + ptratio + lstat + medv , data = rawTrain)
model0ne
```

```
##
## Call:  glm(formula = target ~ zn + indus + chas + nox + rm + age + dis +
##       rad + ptratio + lstat + medv, family = "binomial", data = rawTrain)
##
## Coefficients:
## (Intercept)          zn          indus          chas          nox          rm
## -41.17734    -0.07141   -0.11249    1.25335   49.11180   -0.69362
##      age          dis          rad      ptratio      lstat      medv
##   0.03471    0.83505    0.50619    0.38009    0.03387    0.19946
##
## Degrees of Freedom: 465 Total (i.e. Null);  454 Residual
## Null Deviance:      645.9
## Residual Deviance: 196.6    AIC: 220.6
```

```
# squared transformation to age and lstat
```


```
#age before squared
```

```
summary(rawTrain$age)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      2.90  43.88   77.15   68.37   94.10  100.00
```


```
#age before squared
```

```
qplot(rawTrain$age)
```



```
#age after squared
```


```
qplot((rawTrain$age)^(2))
```




```
#lstat before log  
summary(rawTrain$lstat)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   
##   1.730   7.043   11.350   12.631   16.930   37.970
```

```
#lstat before log  
qplot(rawTrain$lstat)
```



```
#lstat after log  
qplot(log10(rawTrain$lstat))
```



```
#remove Tax squared age and log lstat
modelTwo <- glm(target ~ zn + indus + chas + nox + rm + age^2 + dis + rad + ptratio + log10(lstat) + medv, family = "binomial", data = rawTrain)
modelTwo
```

```
##
## Call:  glm(formula = target ~ zn + indus + chas + nox + rm + age^2 +
##       dis + rad + ptratio + log10(lstat) + medv, family = "binomial",
##       data = rawTrain)
##
## Coefficients:
## (Intercept)          zn          indus          chas          nox
## -40.12700    -0.06734   -0.10958     1.33475    49.33564
##          rm          age          dis          rad          ptratio
##  -0.89104     0.03896     0.84436     0.51164     0.39222
## log10(lstat)          medv
##  -0.14881     0.19872
##
## Degrees of Freedom: 465 Total (i.e. Null);  454 Residual
## Null Deviance:      645.9
## Residual Deviance: 197   AIC: 221
```

```
#remove Tax squared age and log lstat - log dis and zn
modelThree <- glm(target ~ log10(zn + 1) + indus + chas + nox + rm + age^2 + log10(dis) + rad + ptratio + medv, family = "binomial", data = rawTrain)
modelThree
```

```
##
## Call: glm(formula = target ~ log10(zn + 1) + indus + chas + nox + rm +
##       age^2 + log10(dis) + rad + ptratio + log10(lstat) + medv,
##       family = "binomial", data = rawTrain)
##
## Coefficients:
##   (Intercept)  log10(zn + 1)      indus      chas      nox
##   -46.69939    -1.00777    -0.07120    1.11180    54.23952
##      rm      age  log10(dis)      rad  ptratio
##   -1.01689    0.04484    10.03136    0.55084    0.41541
## log10(lstat)      medv
##    0.12432    0.23433
##
## Degrees of Freedom: 465 Total (i.e. Null);  454 Residual
## Null Deviance:      645.9
## Residual Deviance: 189.2    AIC: 213.2
```

```
#remove Tax squared age and log lstat - log dis - log zn
modelFour <- glm(target ~ zn + indus + chas + nox + rm + age^2 + log10(dis) + rad + ptratio + log10(lst
modelFour
```

```
##
## Call: glm(formula = target ~ zn + indus + chas + nox + rm + age^2 +
##       log10(dis) + rad + ptratio + log10(lstat) + medv, family = "binomial",
##       data = rawTrain)
##
## Coefficients:
##   (Intercept)      zn      indus      chas      nox
##   -45.71222   -0.04939   -0.06976    1.17911   53.13061
##      rm      age  log10(dis)      rad  ptratio
##   -1.05455    0.04416    9.47828    0.55075    0.44318
## log10(lstat)      medv
##   -0.16944    0.23047
##
## Degrees of Freedom: 465 Total (i.e. Null);  454 Residual
## Null Deviance:      645.9
## Residual Deviance: 189.8    AIC: 213.8
```

NEXT I WANT TO TRY BOX COX TRANSFORMATIONS

WEE NEED QQ PLOTS AND ACCURACY